

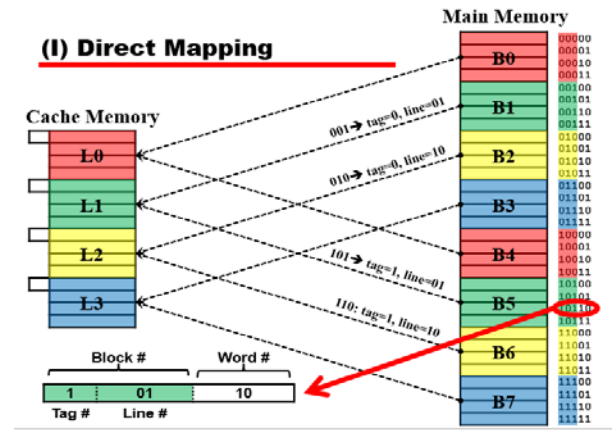
Cache Memory:

$$T_{av} = T_c + (1 - H) T_m$$

1. Mapping:

(I) Direct Mapping

- Each MM block maps to only one cache line.
- Block #**b** maps to line #(**b modulo m**).
- If a block is in cache, it must be in one specific place.




Address Format:

- Memory address is split (based on block size) into:
 - Least significant **w** bits → identify **word** in block.
 - $w = \log_2$ (# of words in block).
 - Most significant **s** bits → identify **block** in MM.
 - $s = \log_2$ (# of blocks in MM).
 - The **s** bits are split (based on cache size) into:
 - Least significant **r** bits → identify cache **line**.
 - $r = \log_2$ (# of lines in cache).
 - Most significant **s - r** bits → **tag** (identify group).

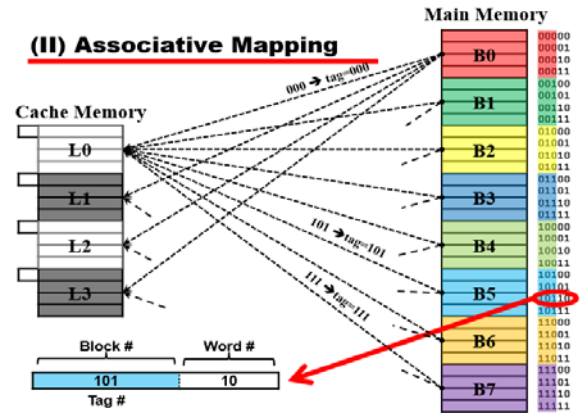
Address Format Summary:

- Address length = $(s + w)$ bits.
- Number of addressable units = 2^{s+w} words.
- Block size = line size = 2^w words.
 - $w = \log_2$ (# of words in block).
- Number of blocks in MM = $M = 2^{s+w}/2^w = 2^s$.
 - $s = \log_2$ (# of words in MM / # of words in block)
- Number of lines in cache = $m = 2^r$.
 - $r = \log_2$ (# of words in cache / # of words in line).
- Size of tag = $(s - r)$ bits.
 - $(s - r) = \log_2$ (# of words in MM / # of words in cache)



(II) Fully-Associative Mapping:

- A main memory block can load into any line of cache.
- Memory address is interpreted as tag and word.
- Tag uniquely identifies block of memory.
- Every line's tag is examined for a match.
- Cache searching gets expensive.




Address Format:

- Memory address is split (based on block size) into:
 1. Least significant **w** bits → identify **word** in block.
 - $w = \log_2 (\text{\# of words in block})$.
 2. Most significant **s** bits → identify **block** in MM, and used as a **tag**.
 - $s = \log_2 (\text{\# of blocks in MM})$.
- In the “tiny example”: $w=2, s=3$.

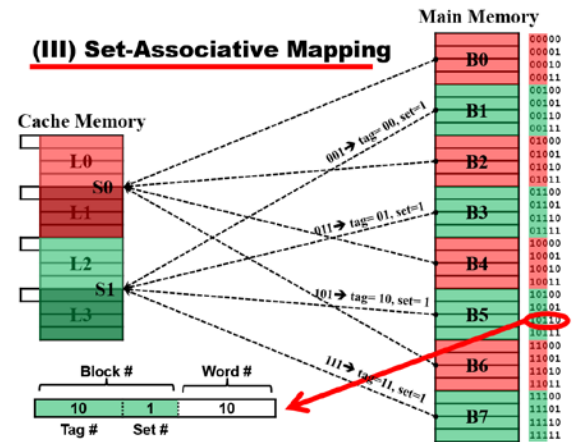
Address Format Summary:

- Address length = $(s + w)$ bits.
- Number of addressable units = 2^{s+w} words.
- Block size = line size = 2^w words.
 - $w = \log_2 (\text{\# of words in block})$.
- Number of blocks in MM = $M = 2^{s+w} / 2^w = 2^s$.
 - $s = \log_2 (\text{\# of words in MM} / \text{\# of words in block})$
- Number of lines in cache = m → unknown!
- Size of tag = s .



(III) Set-Associative Mapping:

- Cache is divided into a number of **sets** (v) of equal size.
- Each set contains a number of lines (k).
 - **k-way** set-associative mapping!
- A block b could map to **any line** in a set i if and only if $i = b \text{ modulo } v$.




Address Format:

- Memory address is split (based on block size) into:
 - Least significant w bits ➔ identify **word** in block.
 - $w = \log_2 (\# \text{ of words in block})$.
 - Most significant s bits ➔ identify **block** in MM.
 - $s = \log_2 (\# \text{ of blocks in MM})$.
 - The s bits are split (based on cache size) into:
 - Least significant d bits ➔ identify cache **set**.
 - $d = \log_2 (\# \text{ of sets in cache})$.
 - Most significant $s - d$ bits ➔ **tag** (identify group).
- In the “tiny example”: $w=2, s=3, d=1, s-d=2$.

Address Format Summary:

- Address length = $(s + w)$ bits.
- Number of addressable units = 2^{s+w} words.
- Block size = line size = 2^w words.
 - $w = \log_2 (\# \text{ of words in block})$.
- Number of blocks in MM = $M = 2^{s+w} / 2^w = 2^s$.
 - $s = \log_2 (\# \text{ of words in MM} / \# \text{ of words in block})$
- Number of lines in set = k ➔ **k-way set-associative**
- Number of sets = $v = 2^d$.
- Number of lines in cache = $m = k * v = k * 2^d$.
 - $d = \log_2 (\# \text{ of words in cache} / \# \text{ of words in set})$.
- Size of tag = $(s - d)$ bits.
 - $(s-d) = \log_2 (\# \text{ of blocks in MM} / \# \text{ of sets in cache})$.



2. Replacement Algorithms:

Direct mapping:

- No choice.
- Each block only maps to one line.
- Replace that line.

Associative and Set Associative:

- Hardware implemented algorithms!
- Most common ones:
 - 1. Least Recently Used (LRU)**
 - Replace the block that has not been recently accessed.
 - 2. First In First Out (FIFO)**
 - Replace the oldest block.
 - 3. Least Frequently Used**
 - Replace the block which has had the fewest hits.
 - 4. Random**
 - Replace any block!

3. Read/Write Policies:

Read Policy (Hit & Miss):

- **Reads dominate processor cache accesses.**
 - All instructions need to be fetched → reads.
 - Some instructions may read one or more operands.
 - Most instructions do not write to memory!

- **What if CPU wants to read a word from MM?**
 - **Read-hit policy:** if read hits in cache (i.e., container block is in cache), just read word from cache!

 - **Read-miss policy:** if read misses in cache (i.e., container block is in not cache),
 1. Read-through: read word from MM, and bring container block to cache.
 2. No-read-through: bring container block to cache, and then read word from cache.

Write Policy (Hit):

- What if CPU wants to write a word to memory?
 - **Write-hit policy:** if write hits (i.e., container block is in cache)
 1. **Write-through:** write word to both cache and MM
 - + **Pros:** easier to implement, MM is always consistent with cache, read miss never results in writes to MM.
 - + **Cons:** Every write needs a MM access, slower writes, more MM bandwidth consumption.
 2. **Write-back:** write word to cache only, do not write word to MM until cached version of container block is replaced.
 - + In addition to tag bits, each cache line stores a “dirty bit”.
 - + **Pros:** multiple writes within block require only one write to MM, faster writes, less MM bandwidth consumption.
 - + **Cons:** harder to implement, MM is not always consistent with cache, reads may cause writes of dirty blocks to MM.

Write Policy (Miss):

- **Write-miss policy:** if write misses (i.e., container block is in not cache)
 1. **Write-allocate:** bring container block to cache, and then act according to write-hit policy.
 2. **No-write-allocate:** write word to container block in MM, and do not bring block to cache.
- **Note:** either write-miss policy could be used with write-through or write-back, however:
 - Write-back caches generally use write-allocate
 - Hoping that subsequent writes to block get captured by cache
 - Write-through caches often use no-write-allocate
 - since subsequent writes block will still have to go to MM